# Snowflake: Attaining Data-Driven Insights with Snowflake and Microsoft Power BI

Alan Biju Palayil

Department of Information Technology, University of the Cumberlands

MSDS-535: Data Mining (MSDS-535-A01)

Professor: Dr. Kathryn Wifvat

Date: May 14th, 2026

# Attaining Data-Driven Insights with Snowflake and Microsoft Power BI

## Introduction

Modern organizations generate enormous amounts of enterprise data from operational systems, cloud environments, applications, customer interactions, and business transactions. To convert this information into actionable business intelligence, organizations need scalable data warehouses, analytical processing systems and visualization platforms. Snowflake and Microsoft Power BI are modern, cloud-native ecosystems that can support large-scale data warehousing, online analytical processing (OLAP) and business intelligence workflows.

The Snowflake and Microsoft Power BI workshop illustrated how the combination of cloud-based data warehousing and interactive analytics and reporting platforms can enable organizations to support enterprise decision-making. Snowflake provides scalable storage, SQL analytics and cloud compute resources. Power BI is a tool for enterprise data visualization, dashboarding and interactive exploration.

In the report, we discuss how to build databases in Snowflake from the existing large-scale datasets, analyze the data warehouse modeling and OLAP operations, and whether to recommend the methods of data cube computation to build business intelligence databases in Snowflake.

## How to build Snowflake databases from large existing data sets

Snowflake is a cloud native architecture that makes it easy to build and run enterprise databases from structured and semi-structured data sources. Organizations can also ingest existing databases, cloud storage systems, data lakes and operational systems directly into Snowflake via automated ingestion pipelines and scalable warehouse environments.

Building a database in Snowflake usually involves creating warehouses, schemas, databases and user roles. Then, with SQL-based ingestion pipelines, external stages and ETL workflows, existing datasets in cloud storage systems such as Amazon S3, Azure Blob Storage or Google Cloud Storage can be loaded into Snowflake tables. Snowflake enables batch ingestion and real-time streaming pipelines that allow processing of enterprise data at scale.

Snowflake's architecture separates storage and compute resources, which makes it very well suited to work with big data. This gives organizations the flexibility to scale warehouses independently of the needs of analytical workloads. Many users and business intelligence tools can access the same datasets stored centrally with little effect on performance concurrently.

The joint solution also improves enterprise analytics, enabling users to connect dashboards directly to Snowflake warehouses, adding to Snowflake's integration with Microsoft Power BI. This architecture allows business users to create reports, visualize trends and analyze enterprise metrics in near real time.

Current data warehouses (Han et al., 2022) are defined as the integration of data from multiple heterogeneous sources to support enterprise decision making and analytical processing. Snowflake is an example of this, with a cloud-native approach to scaling, centralized storage, and SQL-based analytical workflows.

# Data Warehouse Modeling, OLAP Operations, and Data Cube Computation

Data warehouse modeling is the key activity of designing enterprise data for best analytics and reporting. Data warehouse models include star schemas, snowflake schemas, and fact constellation schemas. In a star schema, fact tables are at the center and linked to dimension tables, while in a snowflake schema, dimensions are normalized into several related tables for more consistency and less redundancy.

These schemas can be leveraged in Snowflake environments to enable business intelligence workflows, financial reporting, customer analytics and operational dashboards. Fact tables hold the measurable metrics of the business such as totals of sales, transaction amounts or KPIs of operations. Dimension tables contain descriptive attributes such as customer information, geographic regions or product categories.

Online Analytical Processing (OLAP) operations are an integral part of enterprise analytics systems. OLAP performs roll-up, drill-down, slice, dice and pivot operations to provide multidimensional analysis of data in the warehouse. Roll-up operations give aggregate views of data and drill-down operations enable analysts to drill down to the detailed transactional records. Slice and dice operations are used to filter multi-dimensional data sets, and to pivot analytical perspectives across dimensions.

Snowflake and Power BI are a good combination for OLAP-style analytics. Snowflake warehouses do the execution and aggregation of large-scale queries. Power BI provides interactive dashboards and visualizations. This architecture enables organizations to perform exploratory analysis, trend identification and KPI monitoring efficiently.

Another important concept in OLAP systems and data warehousing is data cube computation. Data cubes are multi-dimensional structures for data analytical processing and data summarization. Data cube computation techniques enhance query performance by pre-aggregating and pre-summarizing data along dimensions such as time, location, products or customers.

The content of the lectures includes different approaches to the data cube computation: full materialization, partial materialization, iceberg cubes and multiway aggregation techniques. With these methods, organizations can optimize their storage utilization, query speed and analytical performance based on their business needs and data volume.

## Recommendation for Data Cube Computation Methods in Snowflake

For building business intelligence databases in Snowflake, especially for enterprise analytics scale, I would recommend using data cube computation techniques. Query performance is improved with data cube computation and multidimensional analysis over large datasets is faster.

But the actual calculation method should be determined according to the organizational needs and workload characteristics. Full materialization allows fast queries, but can be expensive and storage intensive for very large data sets. Partial materialization and iceberg cube computation approaches are typically more efficient as they only store important aggregations or high value analytical combinations.

The scalable warehouse architecture of Snowflake is well suited for implementing optimized data cube strategies. The ability to scale compute resources allows organizations to balance storage costs with analytical performance. Snowflake with Power BI allows organizations to create responsive dashboards and interactive analytical solutions and benefits from the scalability of the cloud-native.

Star schemas may be more suited to generating reports and dashboards for business users, while the snowflake schema may be more appropriate for highly normalized enterprise environments. The organization will have to decide on modeling approaches based on the complexity of the reporting, the performance requirements, and the analytical workload demands.

Multidimensional analytical systems and data warehouse optimization techniques are essential in converting enterprise data into useful business intelligence (Aggarwal, 2015). Snowflake's architecture aligns well with these objectives, offering cloud scalability, SQL analytics, and business intelligence integrations.

# Conclusion

The Snowflake and Microsoft Power BI workshop demonstrated how modern cloud-native architectures are empowering enterprise data warehousing, OLAP operations and business intelligence workflows. Power BI provides you with enterprise analytics, interactive dashboards and visualization tools, while Snowflake provides you with scalable storage, SQL analytics and cloud warehouse management.

Snowflake makes it easy to build databases with large enterprise datasets by providing cloud ingestion pipelines, scalable warehouses, and centralized storage systems. These are data warehouse modeling techniques like star schema and snowflake schema. Used to efficiently provide analytical processing. OLAP operations allow users to view enterprise data from a multidimensional perspective.

The methods of data cube computation are highly recommended in the business intelligence environment as they improve the query performance and support multidimensional analysis over large datasets. Together with Power BI, Snowflake offers a scalable and effective platform for modern enterprise analytics and data-driven decision making for organizations.

# References


Aggarwal, C. C. (2015). Data mining: The textbook. Springer.

Han, J., Pei, J., & Kamber, M. (2022). Data mining: Concepts and techniques (4th ed.). Elsevier.

Kimball, R., & Ross, M. (2013). The data warehouse toolkit: The definitive guide to dimensional modeling (3rd ed.). Wiley.