

AWS Dev Day: Building Modern Machine Learning Applications on Snowflake

Alan Biju Palayil

Department of Information Technology, University of the Cumberlands

MSDS-535: Data Mining

Professor Name

Date: May 18th, 2026

AWS Dev Day: Building Modern Machine Learning Applications on Snowflake

Introduction

Today's organizations generate large volumes of structured and semi-structured data from operational systems, cloud-based applications, IoT devices, financial systems, and customer interactions. To make this information actionable, organizations are increasingly leveraging cloud-native machine learning platforms that can power scalable analytics, data engineering and predictive modeling workflows. Snowflake and Amazon SageMaker offer a modern ecosystem for building machine learning applications with a centralized cloud data warehouse and scalable machine learning infrastructure.

The AWS Dev Day workshop showed how Snowflake and Amazon SageMaker can be integrated to support enterprise-scale machine learning pipelines. Snowflake is a scalable cloud data warehouse, while SageMaker provides machine learning development, model training, and model deployment capabilities. The combination of these technologies provides organizations with the capability to ingest data, pre-process datasets, build predictive models, and persist predictions for enterprise analytics and business intelligence workflows.

In this report, we explain how data can be ingested into Snowflake for analytics, how SageMaker Notebook instances can be configured to have Snowflake connectivity, how data can be loaded into Pandas Data Frames, how machine learning models can be trained using Amazon SageMaker, and how predictions can be stored in Snowflake for evaluation and analytical processing.

Ingesting Data for Analytics Powered by Snowflake

Data ingestion is one of the most crucial components of modern machine learning pipelines. Snowflake provides ingestion support for structured, semi-structured and streaming data from multiple enterprise systems including relational databases, cloud storage platforms, APIs and real-time event streams.

Organizations can load data into Snowflake from cloud storage services like Amazon S3, Azure Blob Storage or Google Cloud Storage. Snow pipe allows for automated continuous ingestion of files as they become available in cloud storage locations. Batch ETL pipelines are also built with SQL scripts, orchestration tools, or native cloud data integration services.

Snowflake is very effective for analytics due to its architecture which separates storage from compute resources. This allows organizations to dynamically scale warehouses based on workload requirements, while maintaining centralized storage. You can run multiple analytical workloads and machine learning processes simultaneously without a significant performance hit.

The workshop showed how enterprise datasets can be loaded into Snowflake warehouses and prepared for downstream analytics and machine learning workflows. Han et al. (2022) state that advanced analytics and large-scale data mining applications require scalable data warehousing systems. Snowflake's architecture is designed to meet these goals, with efficient cloud-native storage and analytical processing.

Launching a SageMaker Notebook Instance with the Snowflake Python Connector

With Amazon SageMaker, you get a managed notebook environment for machine learning development, experimentation, and deployment. SageMaker Notebook instances were provisioned during the workshop, and the Snowflake Python connector was installed on these instances to allow secure communication between Snowflake warehouses and machine learning workflows.

The Snowflake Python connector enables developers and data scientists to run SQL queries and pull data from the warehouse, as well as integrate Snowflake datasets directly into Python-based machine learning environments. SageMaker makes this easy. Notebook instances are provided with preconfigured environments that include all the Python libraries, machine learning frameworks, and scalable compute resources you'll need.

You can configure secure authentication methods like user names and passwords, OAuth integrations, and key-based authentication to establish connections between SageMaker and Snowflake environments. Organizations can connect and leverage enterprise datasets stored in Snowflake warehouses without large data movements between systems.

The combination of Snowflake and SageMaker reduces operational complexity by enabling organizations to consolidate their enterprise data and use scalable machine learning infrastructure in AWS. This architecture supports collaborative workflows between data engineers, analysts, and data scientists.

Building Connections and Pulling Data into Pandas Data Frames

Once a connection is made between SageMaker and Snowflake, you can run SQL queries directly against Snowflake warehouses to pull datasets for machine learning analysis. The Snowflake Python connector can load query results into Pandas Data Frames, which are widely used in Python-based analytics and machine learning workflows.

Pandas Data Frames offer an easy way to preprocess, engineer features, filter, aggregate and explore data. Data scientists can also clean missing values, normalize variables, encode categorical attributes and prepare datasets for training machine learning models.

The workshop showed how SQL queries can be used to extract enterprise data from Snowflake warehouses and the results can be loaded into Python environments for downstream analysis. This workflow takes advantage of the scalability of Snowflake's cloud warehouse architecture and the flexibility of Python analytics libraries such as pandas, NumPy, and scikit-learn.

Preprocessing is a critical step because machine learning models need high quality structured data to make reliable predictions. Feature engineering and preprocessing play a major role in model performance and predictive accuracy, as Géron (2022) points out. With the integration of Snowflake with Pandas and SageMaker, organizations can benefit from scalable and flexible environments for efficient machine learning dataset preparation.

Training Machine Learning Models Using Amazon SageMaker

Amazon SageMaker offers scalable infrastructure to train machine learning models on enterprise data sets from Snowflake warehouses. The workshop showed how to create machine learning workflows using Python libraries and SageMaker notebook environments.

It supports a rich set of machine learning techniques including classification, regression, clustering and predictive analytics models that can be deployed in SageMaker. Common classification algorithms in enterprise analytics applications include decision trees, logistic regression, and ensemble methods, and are used for fraud detection, customer segmentation, cybersecurity analytics, and operational forecasting.

Decision tree classifiers are of particular interest because of their interpretability and capability of supporting multidimensional data analysis. Another common method for binary classification problems and probability estimation is logistic regression. These notions of classification are very much in line with the course material on classification methods and predictive modeling.

In machine learning workflows the data is usually split into training and testing subsets, model accuracy is evaluated and parameters are tuned to improve predictive performance. Organizations have evaluation metrics like accuracy, precision, recall, sensitivity, specificity, and ROC curves to measure the effectiveness of a model.

SageMaker provides easy-to-use training with scalable compute infrastructure, built-in algorithms, and managed machine learning services. This allows organizations to efficiently train large scale predictive models, leveraging centralized enterprise data stored in Snowflake.

Persisting Predictions in Snowflake for Analytics and Evaluation

Once machine learning models have generated predictions, the results can be persisted back into Snowflake tables for reporting, analytics, and operational decision-making. Since predictions live in Snowflake, organizations can now incorporate the results of predictive analytics into dashboards, BI tools, and enterprise reporting workflows.

Prediction tables may contain model outputs such as classification labels, probabilities of fraud, scores of customer risk, recommendations, or forecasted values. Machine learning predictions and analytical metrics can be visualized by business intelligence tools such as Microsoft Power BI and Tableau that can connect directly to Snowflake warehouses.

Predictions in Snowflake means better governance, auditability and reproducibility. Organizations can monitor model output, track predictions over time and feed machine learning results into enterprise analytics pipelines. This architecture enables a centralized governance, while demanding less scattered external data stores.

According to Aggarwal (2015), modern machine learning systems need scalable infrastructures to enable predictive modelling and analytical evaluation. Together, Snowflake and SageMaker provide a cloud-native ecosystem to enable these requirements with scalable storage, machine learning integration, and enterprise analytics workflows.

Conclusion

During the AWS Dev Day workshop, the link between Snowflake and Amazon SageMaker was shown to develop scalable machine learning applications that leverage the power of centralized cloud-native architectures. Snowflake provides scalable data warehousing and

analytical processing capabilities and SageMaker supports machine learning development, model training and deployment workflows.

Using cloud-native ingestion pipelines and scalable warehouse architectures, organizations can load enterprise datasets into Snowflake. SageMaker Notebook instances with the Snowflake Python connector allow developers to pull data into Pandas Data Frames for preprocessing and feature engineering. Training of machine learning models can then be performed on scalable SageMaker infrastructure and evaluated using state-of-the-art classification and predictive analytics techniques.

Pulling predictions into Snowflake enhances the integration of enterprise analytics, governance and business intelligence. Snowflake and SageMaker together offer organizations a modern platform for scalable analytics, machine learning and cloud-native data-driven decision-making.

References

Aggarwal, C. C. (2015). **Data mining: The textbook**. Springer.

Géron, A. (2022). **Hands-on machine learning with Scikit-Learn, Keras, and TensorFlow** (3rd ed.). O'Reilly Media.

Han, J., Pei, J., & Kamber, M. (2022). **Data mining: Concepts and techniques** (4th ed.). Elsevier.