

Configuring Virtual Warehouse Compute Clusters and Qlik Replicate Integration in Snowflake

Alan Biju Palayil

Department of Information Technology, University of the Cumberland

MSDS-535: Data Mining

Professor: Dr. Kathryn Wifvat

Date: June 8th, 2026

Configuring Virtual Warehouse Compute Clusters and Qlik Replicate Integration in Snowflake

Introduction

Cloud-native data platforms continue to transform the way enterprises ingest, manage and analyze enterprise data. Snowflake has become a popular solution because it separates storage and compute, delivering scalable analytics and data integration capabilities. One of the key benefits of Snowflake is its **virtual warehouse architecture**, which allows organizations to scale compute independently based on the demand of the workload.

In the Snowflake lab on setting up virtual warehouses and integrating **Qlik Replicate**, we learned how organizations can build modern data pipelines by configuring compute clusters, defining source and target endpoints, staging data, applying transformations, and running replication jobs. These capabilities are useful in environments such as data warehousing, real-time ingestion, analytics and business intelligence.

This report outlines the steps to configure Snowflake virtual warehouse compute clusters and use Qlik Replicate to build and run a replication workflow to Snowflake.

Configuring Virtual Warehouse Compute Clusters

Snowflake uses a virtual warehouse as the compute engine to execute queries, load data, and perform other analytical operations. Organizations can scale compute resources up or down depending on workload demands because virtual warehouses are storage independent.

Users create a warehouse through the Snowflake interface or SQL commands. To set up a virtual warehouse During configuration many settings can be defined:

- warehouse name
- size (X-Small to 6X-Large)
- auto timer suspends
- auto-resume behavior
- minimum and maximum number of clusters
- scale the policy up

The amount of compute power provisioned is dictated by the size of the warehouse. Larger warehouses have higher processing capacity but cost more credits. Multi-cluster warehouses can scale automatically to handle concurrent workloads.

The auto-suspend is one of the most helpful features, it shuts down the warehouse when not in use, helping control costs. **Auto-resume** enables the warehouse to automatically resume when a query is issued.

According to Han et al. (2022) the increasing volumes of enterprise data require scalable compute environments to support modern analytical workloads. This is supported natively by Snowflake's warehouse architecture through elastic compute scale.

How to connect Partner Connect and Qlik Replicate

With Snowflake's **Partner Connect** you have simple, direct access to ecosystem partners and third-party data tools without having to manually set up complex configurations.

Qlik Replicate can be started directly from Snowflake UI and a linked environment can be created between the two platforms with Partner Connect. The sharing of relevant account connection information with the partner tool is automated by Snowflake, which simplifies authentication and configuration.

Qlik Replicate: Near real time data movement between databases, applications and Snowflake environments This makes it useful for ETL pipelines, replication workflows, migration projects and enterprise data integration.

Partner Connect significantly cuts setup time over traditional connector-based integrations. ##

Creating the Source and Target Endpoints in Qlik Replicate

Once in Qlik Replicate, the next step is to configure **source and target endpoints**.

Source Endpoint

The source endpoint defines where data originates. Examples include:

- relational databases
- transactional systems
- operational data stores
- cloud databases

- file-based sources

Source configuration typically includes:

- host/server name
- database name
- credentials
- connection settings

Target Endpoint

The target endpoint defines where the replicated data will be delivered. In this case, Snowflake acts as the target.

Target configuration typically includes:

- Snowflake account URL
- warehouse
- database
- schema
- authentication credentials

Once both endpoints are configured successfully, Qlik Replicate can begin moving data between them.

Configuring Snowflake Internal or External Staging

Before loading data into Snowflake, staging must be configured. Snowflake supports both internal and external staging.

Internal Staging

Internal stages store data inside Snowflake-managed cloud storage. This is typically easier to configure and manage.

Advantages:

- simpler setup
- native Snowflake management
- good for smaller or internal workloads

External Staging

External stages reference external cloud storage such as:

- Amazon S3
- Azure Blob Storage
- Google Cloud Storage

Advantages:

- integrates with existing data lakes
- scalable for larger enterprise pipelines
- supports broader cloud architecture workflows

The lab demonstrated configuring Qlik Replicate to write data into one of these staging layers before ingestion into Snowflake tables.

Configuring a Replicate Task

Once endpoints and staging are configured, a replication task can be created.

The replication task defines:

- source tables
- target tables
- load mode
- change data capture settings
- transformation rules
- scheduling behavior

Tasks can run as:

- full load
- full load + CDC (change data capture)
- CDC only

This flexibility allows organizations to perform one-time migrations or continuous replication depending on business requirements.

Configuring Transformations and Testing Expressions

One valuable feature demonstrated in the lab was configuring transformations inside Qlik Replicate.

Transformations allow data to be modified before loading into Snowflake. Examples include:

- renaming columns
- changing data types
- filtering records
- creating calculated fields
- string manipulation
- applying conditional expressions

Testing expressions ensures transformation logic behaves correctly before the task runs in production.

This step is important because it improves data quality and ensures Snowflake receives properly structured data for downstream analytics.

Running the Task

After configuration is complete, the replication task can be executed.

When the task runs:

- Qlik Replicate reads source data
- applies configured transformations
- stages the data
- loads it into Snowflake target tables

Users can monitor:

- task progress
- row counts

- replication status
- errors
- performance metrics

Once complete, the data becomes available inside Snowflake for SQL querying, reporting, analytics, or machine learning workflows.

Aggarwal (2015) explains that efficient ingestion and transformation pipelines are critical for converting enterprise data into usable analytical information. Replication tools like Qlik Replicate support this by simplifying movement between systems.

Conclusion

The virtual warehouse and Qlik Replicate lab demonstrated how Snowflake supports scalable compute management and modern data integration workflows.

Configuring virtual warehouses allows organizations to allocate compute resources dynamically while controlling cost through auto-scaling and auto-suspend capabilities. Partner Connect simplifies integration with Qlik Replicate, enabling fast setup for data replication workflows.

Through Qlik Replicate, organizations can create source and target endpoints, configure staging, define replication tasks, apply transformations, and run automated data pipelines into Snowflake.

Together, these capabilities make Snowflake highly effective for enterprise data warehousing, replication, ETL processing, and cloud-based analytics environments. The

flexibility, scalability, and integration options demonstrated in the lab make Snowflake a strong platform for modern enterprise data operations.

References

Aggarwal, C. C. (2015). *Data mining: The textbook*. Springer.

Han, J., Pei, J., & Kamber, M. (2022). *Data mining: Concepts and techniques* (4th ed.). Elsevier.

Kimball, R., & Ross, M. (2013). *The data warehouse toolkit: The definitive guide to dimensional modeling* (3rd ed.). Wiley.